

# Bayesian Multi-level Regression

Bayesian Modeling in brms

---

Michael D. Nunez

29-August-2022

Amsterdam Mathematical Psychology Laboratory  
University of Amsterdam, The Netherlands

# Sections

The data

Regression in brms

ANOVA in brms

Reading brms output

Structured individual differences

Multi-level regression in brms

## The data

---

## Obtaining the real data set

```
# install.packages('curl')
library(curl)

# See https://github.com/mdnunez/encodingN200
pdmdat <- curl("https://tinyurl.com/dataBayesCogMod")
pdm <- read.csv(pdmdat)

head(pdm)
```

## Regression in brms

---

# Simple linear regression equations

- $y_i$  are the observations of our dependent variable  $y$  for each observation  $i$
- $x_i$  are the observations of our independent variable
- $y_i \sim \text{Normal}(\mu_i, \sigma^2)$
- $\mu_i = \beta_0 + \beta_1 x_i$
- You **may** have previously learned this, which is equivalent:
- $y_i = \beta_0 + \beta_1 x_i + \epsilon_i$
- $\epsilon_i \sim \text{Normal}(0, \sigma^2)$

# Simple linear regression using base R

```
# install.packages('brms')  
library(brms)
```

- brms nicely follows the same standard R formula syntax that base R and many packages follow
- Here is an example of the base R code:

```
simple_lm <- lm(RT ~ N200_latencies, data=pdm)  
summary(simple_lm)
```

# Simple linear regression using brms

```
# install.packages('brms')  
library(brms)
```

- brms nicely follows the same standard R formula syntax that base R and many packages follow
- Here is an example of brms code:
- Note that we used the **default Stan prior**

```
bayes_lm <- brm(RT ~ N200_latencies, data=pdm)  
summary(bayes_lm)
```



What is the brms code to estimate a linear regression with *RT* as the dependent variable, *N200\_latencies* and *N200\_amplitudes* as the independent variables, and *an interaction term*?

Hint: Read the help file in RStudio using `?brm`

## Extended linear regression using brms

- Note there are at least three different methods that yield the same solution.
- Here is one **short** example solution:

```
bayes_lmext <- brm(RT ~ N200_latencies*N200_amplitudes,  
  data=pdm)  
summary(bayes_lmext)
```

# Extended linear regression equations

- $y_i$  are the observations of our dependent variable  $y$  for each observation  $i$
- $x_{ki}$  are the observations of our independent variables for each independent variable  $k$
- $y_i \sim \text{Normal}(\mu_i, \sigma^2)$
- $\mu_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \beta_3 x_{3i} + \dots$
- You **may** have previously learned this, which is equivalent:
- $y_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \beta_3 x_{3i} + \dots + \epsilon_i$
- $\epsilon_i \sim \text{Normal}(0, \sigma^2)$

## ANOVA in brms

---

# A simple 3-level, 1-way Bayesian “ANOVA”

- Let's think about the signal-to-noise (SNR) *condition* effect on response time *RT*

```
bayes_anova <- brm(RT ~ factor(condition), data=pdm)
summary(bayes_anova)
```

## A simple 2-way Bayesian “ANOVA”

- Let's think about both the signal-to-noise (SNR) *condition*, the effect of *accuracy*, and the interaction effect on response time (*RT*)

```
bayes_anova <-  
brm(RT ~ factor(condition)*factor(accuracy), data=pdm)  
summary(bayes_anova)
```

## Reading brms output

---

What effects are significant in the previous model?



# What effects are “significant”?

- Use Bayesian probability as evidence for an effect
- For what range are you 95% certain of an effect
- What if this overlaps 0?

## Population-Level Effects:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat
Intercept	0.84	0.01	0.82	0.86	1.00
factorcondition1	-0.05	0.01	-0.07	-0.02	1.00
factorcondition2	-0.02	0.01	-0.04	0.01	1.00
factoraccuracy1	-0.03	0.01	-0.06	-0.01	1.00
factorcondition1:factoraccuracy1	-0.03	0.02	-0.06	0.00	1.00
factorcondition2:factoraccuracy1	-0.06	0.02	-0.09	-0.02	1.00

## Bulk\_ESS Tail\_ESS

Intercept	2108	2570
factorcondition1	2060	2427
factorcondition2	1973	2421
factoraccuracy1	1967	2563
factorcondition1:factoraccuracy1	1943	2339
factorcondition2:factoraccuracy1	1957	2221

## Family Specific Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	0.24	0.00	0.23	0.24	1.00	3513	2710

# Convergence diagnostics: $\hat{R}$

- The *Gelman-Rubin convergence diagnostic* statistic  $\hat{R}$
- $\hat{R}$  give the ratio of the between-chain to within-chain variance.
- $\hat{R}$  be close to 1
- **Rule of thumb:**  $\leq 1.01$  for regression models
- **Rule of thumb:**  $\leq 1.10$  for complex hierarchical models

# Convergence diagnostics: ESS

- The effective sample size **ESS** should be large for appropriate posterior distribution estimates.
- **ESS** statistics penalize the true number of posterior samples by the Markov Chain autocorrelation
- **Bulk\_ESS** is best ESS diagnostic the mean posterior **Estimate**
- **Tail\_ESS** is best ESS diagnostic the 95% credible intervals **CI**
- **Rule of thumb:**  $> 100$

# Did the model converge?

## Population-Level Effects:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat
Intercept	0.84	0.01	0.82	0.86	1.00
factorcondition1	-0.05	0.01	-0.07	-0.02	1.00
factorcondition2	-0.02	0.01	-0.04	0.01	1.00
factoraccuracy1	-0.03	0.01	-0.06	-0.01	1.00
factorcondition1:factoraccuracy1	-0.03	0.02	-0.06	0.00	1.00
factorcondition2:factoraccuracy1	-0.06	0.02	-0.09	-0.02	1.00

## Bulk\_ESS Tail\_ESS

Intercept	2108	2570
factorcondition1	2060	2427
factorcondition2	1973	2421
factoraccuracy1	1967	2563
factorcondition1:factoraccuracy1	1943	2339
factorcondition2:factoraccuracy1	1957	2221

## Family Specific Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	0.24	0.00	0.23	0.24	1.00	3513	2710

## Structured individual differences

---

# Simpson's paradox

- Simpson's paradox refers to the fact that the the effects could be reversed on the participant level

## Fitting a regression per participant

- One way to avoid Simpson's paradox is to fit regression models to each participant / individual
- However the sample size per model is greatly reduced.
- Also this method is less robust to contaminant data

# Structured individual differences

- Another method is to assume that the participant-level differences are related
- Knowledge about other participants helps estimate parameters of new participants
- The sample size per participant necessary to estimate the model is much less than the individual model strategy
- Outlier participants are less influential



# Hierarchical linear regression equations

- $y_{ip}$  are the observations of our dependent variable  $y$  for each observation  $i$  and participant  $p$
  - $x_{ip}$  are the observations of our independent variable
  - $\beta_{0p}$  are participant level intercepts
  - $\beta_{1p}$  be the observations of our independent variables
  - Both  $\beta_{0p}$  and  $\beta_{1p}$  come from *hierarchical* distributions
- 
- $y_i \sim \text{Normal}(\mu_i, \sigma^2)$
  - $\mu_i = \beta_{0p} + \beta_{1p}x_{ip}$
  - $\beta_{0p} \sim \text{Normal}(\mu_0, \sigma_0^2)$
  - $\beta_{1p} \sim \text{Normal}(\mu_1, \sigma_1^2)$

## Multi-level regression in brms

---

# Linear regression with random intercepts

```
# install.packages('brms')  
library(brms)
```

- Here is an example of brms code for random intercepts:
- This may take a couple of minutes to run

```
bayes_randint <-  
brm(RT ~ N200_latencies + (1|subject), data=pdm)  
summary(bayes_randint)
```

- Press **STOP** in the top right of your RStudio console to end the model fitting

# Linear regression with random slopes

```
# install.packages('brms')  
library(brms)
```

- Here is an example of brms code for random slopes:
- **DO NOT RUN THIS NOW**
- This will take some time to run

```
bayes_ranffect <-  
brm(RT ~ N200_latencies + (N200_latencies|subject),  
data=pdm)  
summary(bayes_ranffect)
```

What is the brms code to estimate a linear regression with *RT* as the dependent variable, *N200\_latencies* and *N200\_amplitudes* as the independent variables, *an interaction term*, and a random intercept for each *subject*?

# Final hierarchical regression model in brms

- Note there are a few different methods that yield the same solution.
- Here is one **short** example solution:

```
bayes_final <-  
brm(RT ~ (1|subject) + N200_latencies*N200_amplitudes,  
data=pdm)  
summary(bayes_final)
```

- Press **STOP** in the top right of your RStudio console to end the model fitting

Now let's talk about cognitive modeling!